

Power consumption analysis and optimization

ID: RK-KF-YF-177

Release Version: V1.0.1

Release Date: 2021-04-29

Security Level: ☐Top-Secret ☐Secret ☐Internal ☒Public

DISCLAIMER

THIS DOCUMENT IS PROVIDED "AS IS". ROCKCHIP ELECTRONICS CO., LTD.("ROCKCHIP")DOES NOT PROVIDE ANY WARRANTY OF ANY KIND, EXPRESSED, IMPLIED OR OTHERWISE, WITH RESPECT TO THE ACCURACY, RELIABILITY, COMPLETENESS, MERCHANTABILITY, FITNESS FOR ANY PARTICULAR PURPOSE OR NON-INFRINGEMENT OF ANY REPRESENTATION, INFORMATION AND CONTENT IN THIS DOCUMENT. THIS DOCUMENT IS FOR REFERENCE ONLY. THIS DOCUMENT MAY BE UPDATED OR CHANGED WITHOUT ANY NOTICE AT ANY TIME DUE TO THE UPGRADES OF THE PRODUCT OR ANY OTHER REASONS.

Trademark Statement

"Rockchip", "瑞芯微", "瑞芯" shall be Rockchip's registered trademarks and owned by Rockchip. All the other trademarks or registered trademarks mentioned in this document shall be owned by their respective owners.

All rights reserved. ©2021. Rockchip Electronics Co., Ltd.

Beyond the scope of fair use, neither any entity nor individual shall extract, copy, or distribute this document in any form in whole or in part without the written approval of Rockchip.

Rockchip Electronics Co., Ltd.

No.18 Building, A District, No.89, software Boulevard Fuzhou, Fujian, PRC

Website: www.rock-chips.com

Customer service Tel: +86-4007-700-590

Customer service Fax: +86-591-83951833

Customer service e-Mail: fae@rock-chips.com

Preface

Overview

This document mainly describes some basic concepts and optimization methods of power consumption for RK platform chips.

Product version

Product name	Kernel version
All chips	All kernel versions

Applicable object

This document (the guide) is mainly suitable for below engineers:

Field application engineers

Software development engineers

Revision history

Version	Author	Date	Change Description
V1.0.0	Tim Chen	2019-08-31	Initial version
V1.0.1	Karen Huang	2021-04-29	Modify format and add copyright information

Contents

Power consumption analysis and optimization

1. Basic concept
 - 1.1 Frequency (clk) and voltage
 - 1.2 Voltage domain(VD) and power domain(PD)
 - 1.3 DCDC (Direct Current) and LDO (Low dropout regulator)
 - 1.4 Static power consumption and dynamic power consumption
 - 1.5 DVFS(Dynamic Voltage and Frequency Scaling), CPUFREQ and DEVFREQ
2. Power consumption measurement
 - 2.1 Measurement method
 - 2.2 Measurement tool
3. Power consumption data analysis
 - 3.1 Calculate theoretical power consumption
 - 3.2 Compare with EVB data
 - 3.3 Data analysis for each path
 - 3.3.1 VDD_CORE/VDD_CPU/VDD_ARM
 - 3.3.2 VDD_GPU
 - 3.3.3 VDD_LOGIC
 - 3.3.4 VCC_DDR
 - 3.3.5 VCC_IO
 - 3.4 Common scenario analysis
 - 3.4.1 Static desktop
 - 3.4.2 Video playback
 - 3.4.3 Game
 - 3.4.4 Deepsleep
4. Power consumption optimization strategy
 - 4.1 CPU optimization
 - 4.2 DDR optimization
 - 4.3 Thermal control optimization
 - 4.4 Power optimization

1. Basic concept

1.1 Frequency (clk) and voltage

Generally there are many modules inside SoC, such as ARM, GPU, DDR, I2C, SPI, USB and so on. When each module is working, the digital logic part requires an appropriate frequency and corresponding voltage. The higher the module frequency is, the higher the voltage is required. The frequency and voltage are two important parameters of power consumption.

1.2 Voltage domain(VD) and power domain(PD)

Generally all modules inside SoC have digital logic part and IO part. The digital logic part is mainly responsible for computing and status control, and IO part is mainly responsible for the transmission of the interface signal (some modules don't have IO, such as ARM, GPU, etc.). Generally the power supplies of the digital logic and IO are separated. The power consumption of IO part is generally fixed, while the power consumption of digital logic part changes a lot due to the influence of frequency and voltage. In order to optimize the power consumption, the digital logic part inside the chip is divided into voltage domain and power domain according to the module.

- The voltage domain means the domain where several modules inside the chip share one external power supply. It can adjust or turn on/off the voltage independently. Generally the modules with similar running voltage and not large power consumption can be put in the same voltage domain. But if the power consumption is very large, it is better to use a separate voltage domain, which is convenient to manage the power consumption and also avoid the peak current exceeding the limit of external power supply. To ensure that all modules can work normally, need to set the voltage of the voltage domain to the required voltage of the module with the highest voltage requirement (excluding the closed module).
- One voltage domain may contain many modules and these modules generally don't work at the same time. With power supply, the modules not working will have leakage. In order to reduce the leakage, generally we will divide one voltage domain into several areas, and each area can independently turn on/off the power supply. After some area switches off the power supply, it will be isolated from other modules and significantly reduce the leakage. This kind of area is called power domain.

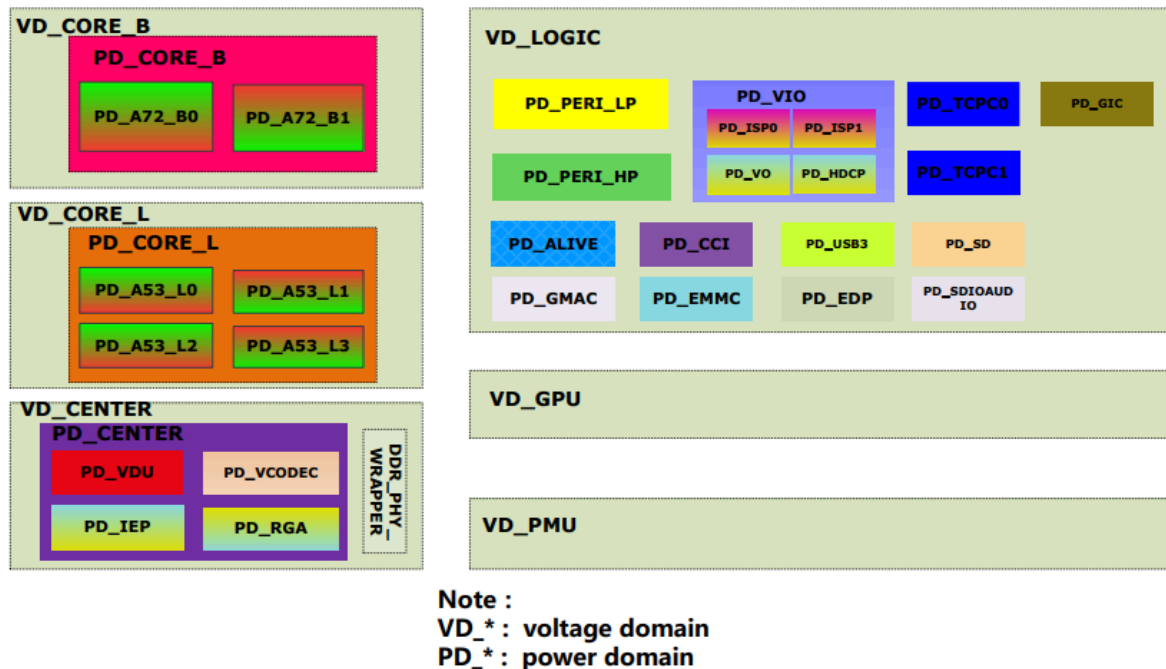
Take RK3399 as example, there are 6 VD:

- VD_CORE_B: including two big cores Cortex-A72, the power consumption is relatively large, so separate a voltage domain.
- VD_CORE_L: including four little cores Cortex-A53, the power consumption is relatively large, so separate a voltage domain.
- VD_LOGIC: including some peripherals' controller and system bus, such as USB, EMMC, GMAC, SPI, I2C, EDP, VOP, AXI, AHB, APB, and so on.
- VD_CENTER: including vdpu, vepu, iep, rga and DDR controller.

- VD_GPU: including GPU, the power consumption is relatively large, so separate a voltage domain.
- VD_PMU: including PMU, SRAM, GPIO, PVTM and other modules relating to suspend and resume process.

The block diagram is as below:

Mclaren power domain & voltage domain



1.3 DCDC (Direct Current) and LDO (Low dropout regulator)

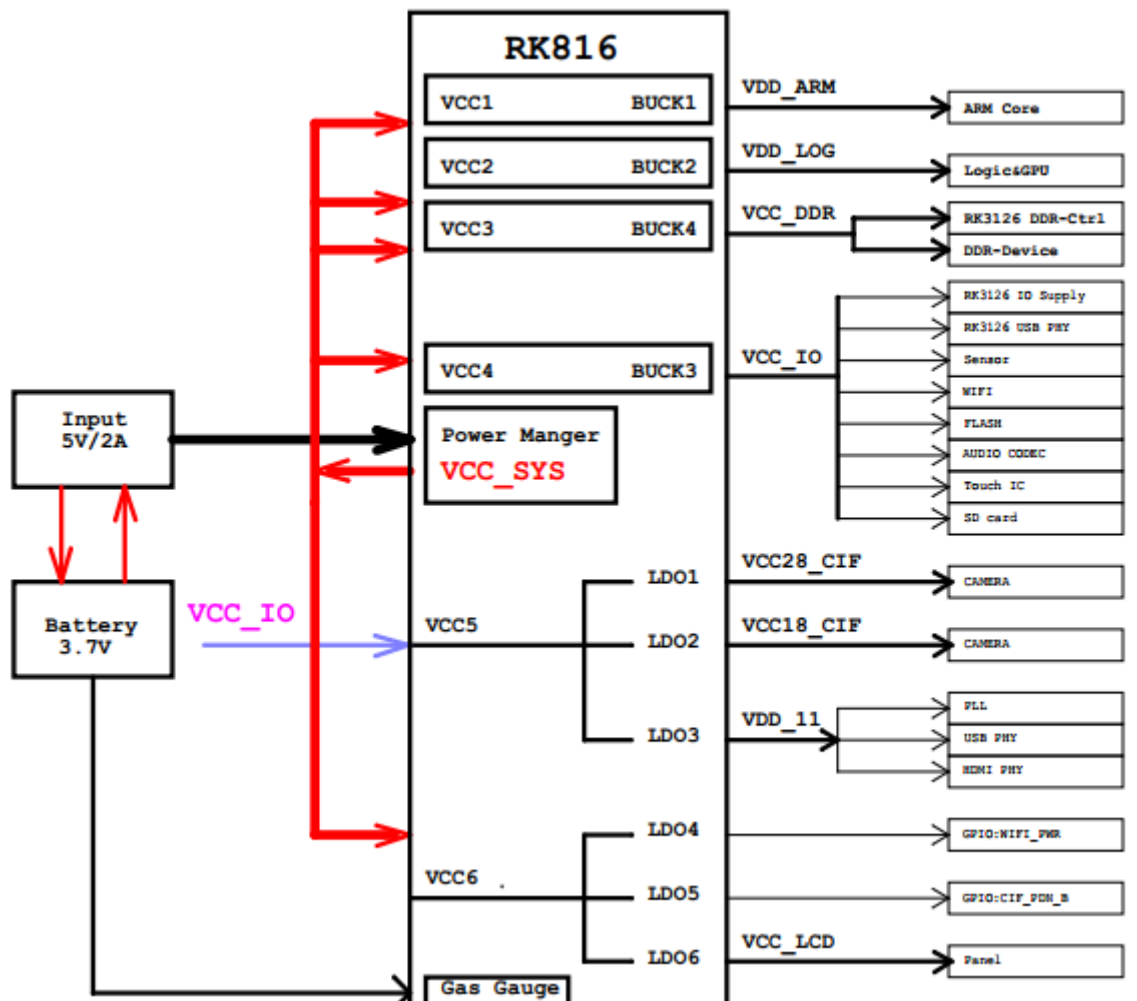
The external power supply of SoC mainly includes DCDC and LDO:

- DCDC generally means switch power, conversion efficiency is high, the efficiency can be up to 80%~90%, when the current is relatively large, need to use DCDC to improve the power efficiency.
- The main characteristic of LDO is, the input current equals to output current, so power efficiency = (output voltage)/(input voltage). Assuming input 3.8V, output 1.0V, the power efficiency is $1V/3.8V=26.3\%$, which indicates this is low efficiency.

Take the power supply solution of RK3126+RK816 as example:

- 4 BUCK of RK816 separately supply power for ARM, LOG, DDR, IO of RK3126, because the current of these modules are all relatively large (BUCK is a kind of voltage drop DCDC).
- 6 LDO of RK816 separately supply power for PLL, PHY and some peripherals of RK3126, because the current of these modules are relatively small.

The block diagram of the power supply is as below:



1.4 Static power consumption and dynamic power consumption

- The static power consumption is the power consumption consumed by the leakage of transistor when the internal modules of SoC are not working. The static power consumption will increase with the increase of the temperature and voltage.
- The dynamic power consumption is the power consumption consumed by the conversion of internal circuit when the internal modules of SoC are working. The dynamic power consumption will increase with the increase of the frequency and voltage.

The format of dynamic power consumption:

/ C is constant, V is voltage, F is frequency*/*

$$P(d) = C * V^2 * F$$

1.5 DVFS(Dynamic Voltage and Frequency Scaling), CPUFREQ and DEVFREQ

The higher the module working frequency and the voltage are, the higher the power consumption is. So need dynamically adjust the frequency and voltage to optimize the power consumption. When the system is idle, reduce the frequency and voltage, when the system is busy, increase the frequency and voltage.

- DVFS is the technology of dynamic voltage and frequency scaling, which is the basic technology implementation of CPUFREQ and DEVFREQ.
- CPUFREQ is the software framework of dynamic CPU frequency scaling, including several different frequency scaling strategies. For more details, please refer to the document 《Rockchip-Developer-Guide-Linux4.4-CPUFreq-CN》.
- DEVFREQ is the software framework of dynamic peripheral(not including CPU) frequency scaling, including several different frequency scaling strategies. For more details, please refer to 《Rockchip-Developer-Guide-Linux4.4-Devfreq》.

2. Power consumption measurement

Before optimizing the power consumption, need to measure the voltage and current of each power supply, analyze the data and then optimize accordingly.

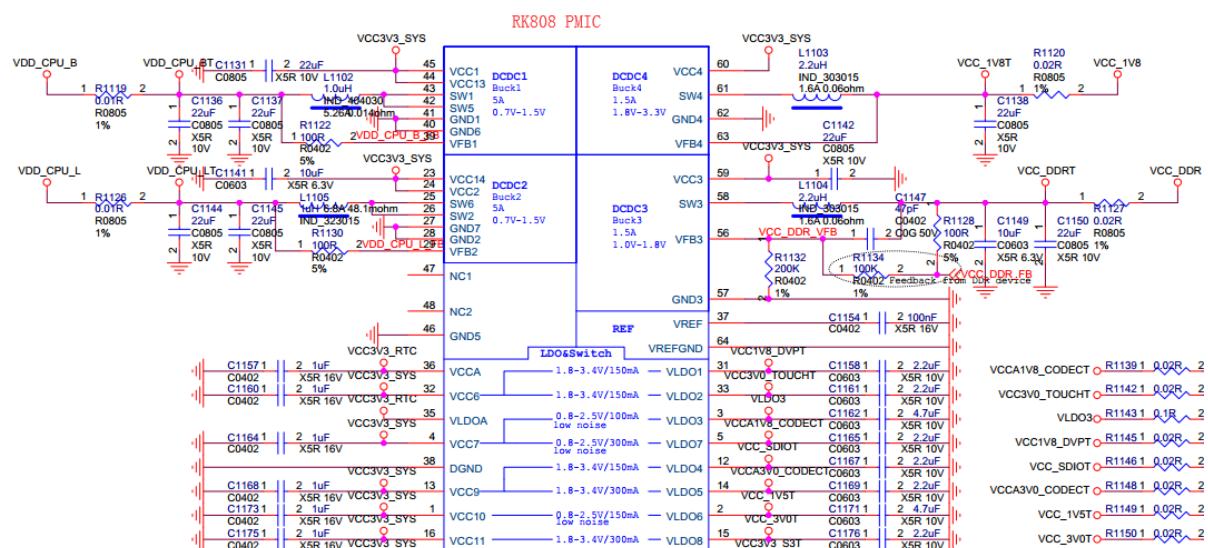
Note: the temperature is an important parameter affecting the power consumption, so need to record the real-time temperature when measuring the power consumption. The command to acquire the temperature is as below:

```
cat /sys/class/thermal/thermal_zone0/temp
```

2.1 Measurement method

Series connect a resistor R in the circuit to measure the voltage difference U between two sides of the resistor, then the current $I=U/R$. Generally here we use the resistor with 0.01 ohm, but you need to adjust the resistance according to the current.

Take RK3399 EVB board as example, by this method, series connect 0.01 ohm resistor to the output of VDD_CPU_B, VDD_CPU_L, VCC_1V8 and VCC_DDR, as shown in below picture:

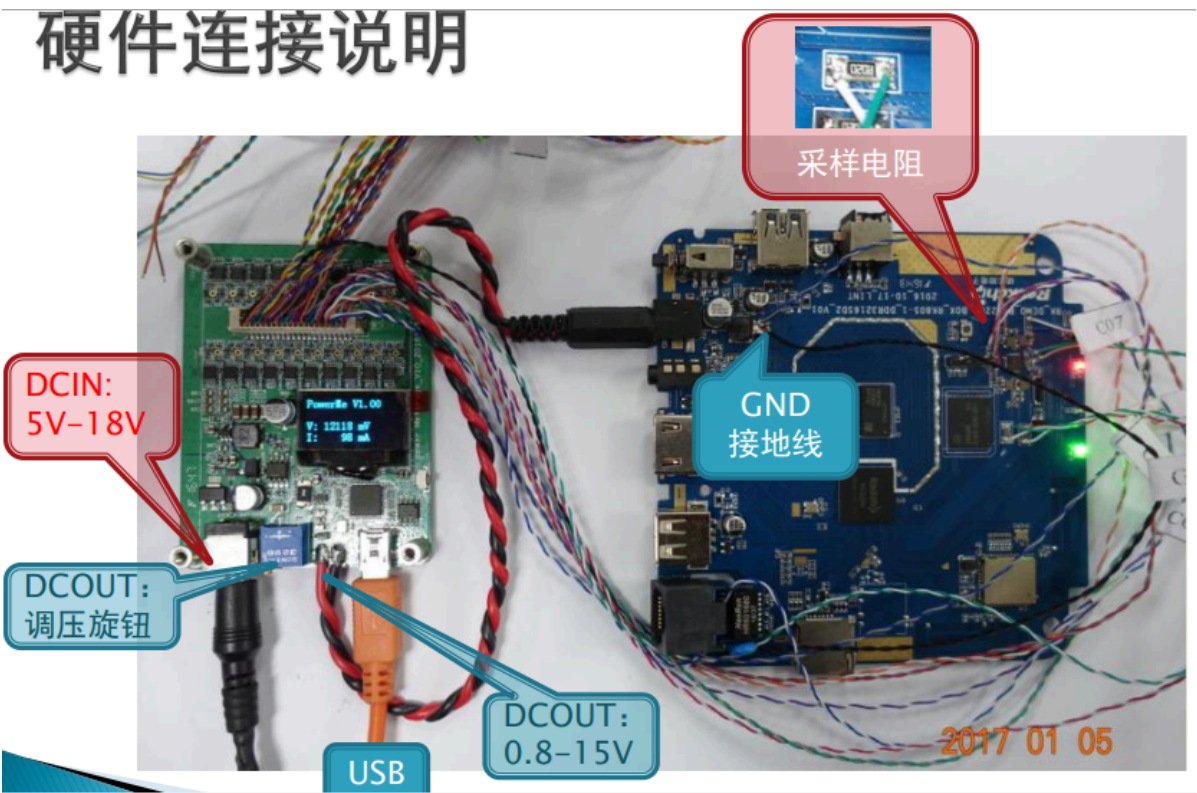


2.2 Measurement tool

As there are many channels of power required to be measured, use multi-channel voltage/current collector can effectively improve the testing efficiency. PowerMeterage is the voltage/current collection tool developed by RockChip and it can measure 20 channels of power consumption data at the same time. The interface is as below:

RK3399 (Sample:13094 1212 Sample/S 00:00:10)																				
Name	12V		VSYS		VCC3V3_SYS		VDD_CPU_B		VDD_LOG		VDD_CPU_L		VDD_CENTER		VCC_DDR		VCC_LVB		VDD_GPU	
Avg	11.891V	258.4mA	5.048V	565.3mA	3.431V	91.1mA	0.812V	39.9mA	0.94V	292.9mA	0.821V	72.3mA	0.909V	433.9mA	1.254V	140.6mA	1.816V	43.8mA	0.798V	199.8mA
Rms	11.891V	260.8mA	5.048V	570.6mA	3.431V	91.1mA	0.814V	130.4mA	0.94V	293mA	0.821V	79.9mA	0.909V	438.8mA	1.254V	147.2mA	1.816V	43.9mA	0.798V	283.5mA
Max	11.92V	523.3mA	5.075V	1156.9mA	3.435V	104.4mA	1.145V	1609.1mA	0.941V	319.6mA	1.074V	423.6mA	0.912V	627.9mA	1.259V	319.4mA	1.819V	70mA	0.802V	559.8mA
Now	11.894V	249.5mA	5.047V	548.9mA	3.431V	91.1mA	0.8V	9.2mA	0.94V	292.2mA	0.817V	67.1mA	0.909V	424.7mA	1.255V	134.5mA	1.816V	43.6mA	0.798V	172.1mA
CH/mA	CH:0	50mA	CH:1	50mA	CH:2	50mA	CH:3	50mA	CH:4	50mA	CH:5	50mA	CH:6	50mA	CH:7	50mA	CH:8	50mA	CH:9	50mA
Name	IR_LED		MIPI_BL		CAMERA_HOST2		VDD_LCD													
Avg	4.963V	0.1mA	18.345V	4.5mA	5.024V	146.9mA	2.787V	20.8mA												
Rms	4.963V	0.1mA	18.345V	4.5mA	5.024V	147.3mA	2.787V	20.8mA												
Max	4.985V	0.2mA	18.36V	4.7mA	5.051V	160.5mA	2.79V	21.9mA												
Now	4.962V	0.1mA	18.344V	4.5mA	5.024V	146.6mA	2.788V	20.8mA												
CH/mA	CH:10	50mA	CH:11	100mA	CH:12	100mA	CH:13	100mA												

The hardware connection of PowerMeterage is as below:



3. Power consumption data analysis

3.1 Calculate theoretical power consumption

Use PowerMeterage tool to break down the power consumption of each path, convert DCDC to the battery with 80% 90% efficiency, the output current of LDO is equal to the input current, convert DCDC, LDO and other powers to the battery, and then add them up to estimate the total power consumption. If it is very different from the power consumption actually measured on the battery, maybe there is leakage. Need to analyze further.

Take RK3326 EVB board as example, the static desktop power consumption is as below:

Note: Because the test result of each path should be converted to the power consumption of the battery, so it is more convenient to compare the actually measured current of the battery with the theoretical current on battery.

Type	power-supply	Voltage(V)	current(mA)	Theoretical current on battery-3.8V(mA)	Remark
DC/DC	VDD_ARM	0.96	10.20	3.23	With 80% efficiency, conversion formula: $V * I / \text{efficiency} / \text{voltage of the battery}$
DC/DC	VDD_LOG	0.96	89.30	28.20	eg: Theoretical current of VDD_LOG on battery(3.8V)= $0.96 * 89.3 / 0.8 / 3.8 = 28.2$
DC/DC	VCC_DDR	1.26	38.50	15.91	
DC/DC	VCC_IO	2.99	4.50	4.43	
LDO	VCC_1V8	1.81	28.80	28.80	Output current of LDO is equal to input current
LDO	VDD_1V0	1.00	10.90	10.90	
LDO	VCC3V0_PMU	3.01	1.20	1.20	
battery	VBAT	3.81	94.60	92.67	Theoretical value is similar to actually measured value

3.2 Compare with EVB data

Break down the power consumption data of each path, compare with the data of EVB in the same scenario, and check if there is problem. For example, the following is the comparison of the static desktop power consumption between RK3326 EVB board and customer device, it can be seen that customer board's power consumption of ARM and LOG are abnormal, and need to analyze further.

Type	power-supply	EVb		Customer device	
		Voltage(V)	Current(mA)	Voltage(V)	Current(mA)
DC/DC	VDD_ARM	0.96	10.20	1.10	212.50
DC/DC	VDD_LOG	0.96	89.30	1.00	151.30
DC/DC	VCC_DDR	1.26	38.50	1.27	40.50
DC/DC	VCC_IO	2.99	4.50	2.99	4.80
LDO	VCC_1V8	1.81	28.80	1.81	29.80
LDO	VDD_1V0	1.00	10.90	1.00	10.20
LDO	VCC3V0_PMU	3.01	1.20	3.01	1.40
battery	VBAT	3.81	94.60	3.81	191.6

3.3 Data analysis for each path

3.3.1 VDD_CORE/VDD_CPU/VDD_ARM

These three names are the same power, that is, ARM core power. This power consumption can be analyzed mainly from the following aspects:

- Confirm if the frequency voltage table (opp-table) is normal or not, if the actually measured voltage is consistent with the set voltage or not.

Relative commands are as below:

```
/* Acquire the frequency voltage table, target column means the voltage required
by some frequency */
cat /sys/kernel/debug/opp/opp_summary
device          rate(Hz)    target(uV)    min(uV)    max(uV)
-----
...
cpu0
           408000000    950000    950000    1350000
           600000000    950000    950000    1350000
           816000000    1000000    1000000    1350000
          1008000000    1125000    1125000    1350000
          1200000000    1275000    1275000    1350000
          1248000000    1300000    1300000    1350000
          1296000000    1350000    1350000    1350000

/* Check the frequency scaling strategy currently used by cpufreq, cpu frequency
scaling is enabled with the default interactive strategy */
cat /sys/devices/system/cpu/cpu0/cpufreq/scaling_governor
interactive
```

```

/* Set userspace strategy to fix the frequency of cpu, then set different
frequencies, compare the set voltage with the measured voltage */
echo userspace > /sys/devices/system/cpu/cpu0/cpufreq/scaling_governor
cat /sys/devices/system/cpu/cpu0/cpufreq/scaling_governor
userspace

/* Check the frequency point supported by cpufreq */
cat /sys/devices/system/cpu/cpu0/cpufreq/scaling_available_frequencies
408000 600000 816000 1008000 1200000 1248000 1296000

/* Set the fixed frequency */
echo 408000 > /sys/devices/system/cpu/cpu0/cpufreq/scaling_setspeed

/* Confirm current frequency*/
cat /sys/devices/system/cpu/cpu0/cpufreq/scaling_cur_freq
408000

/* Confirm current voltage, and compare with measured value, vdd_arm represents
the name of regulator, which is different for different projects*/
cat /sys/kernel/debug/regulator/vdd_arm/voltage
950000

/* Acquire current voltages of all regulators*/
cat /sys/kernel/debug/regulator/regulator_summary
regulator                use open bypass voltage current      min      max
-----
...
vcc3v8_sys                0    12      0 3800mV      0mA 3800mV 3800mV
  deviceless              0mV      0mV
  vdd_logic                0     4      0 950mV      0mA 950mV 1350mV
    dmc                   950mV 1350mV
    ff400000.gpu           950mV 1350mV
    bus-apll               950mV 1350mV
    deviceless             0mV      0mV
  vdd_arm                  0     2      0 950mV      0mA 950mV 1350mV
    cpu0                   950mV 1350mV
    deviceless             0mV      0mV
...

```

- Check cpu loading, analyze if there is irregular task or interrupt.

```

/* Use top command to check the task loading, the output of top with different
versions will have difference, this version of top supports to check the thread
and the running cpu of the thread */
top -m 5 -t
User 51%, System 2%, IOW 0%, IRQ 0%
User 712 + Nice 0 + Sys 33 + Idle 634 + IOW 0 + IRQ 0 + SIRQ 0 = 1379

/* PR column represents currently running cpu of the thread, the sum of all cpu
loading percentage is equal to 100%, so the highest loading percentage of each
cpu is 100%/NR_CPU, the highest loading percentage of each CPU of SoC with 4
cores is 25% */

```

PID	TID	PR	CPU%	S	VSS	RSS	PCY	UID	Thread	Proc
2631	2631	3	25%	R	3104K	552K	fg	root	busybox	busybox
2632	2632	2	25%	R	3104K	552K	fg	root	busybox	busybox

```

2633 2633 1 3% R 740K 400K fg root top /data/top
255 476 0 0% S 15492K 4988K fg system HwBinder:255_1
/vendor/bin/hw/android.hardware.sensors@1.0-service
419 478 1 0% S 3770752K 256884K fg system SensorService system_server

/* Use cpustats to observe the frequency change of cpu */
cpustats
Total: User 600 + Nice 0 + Sys 3 + Idle 591 + IOW 0 + IRQ 0 + SIRQ 0 = 1194
408000kHz 0 +
600000kHz 0 +
816000kHz 0 +
1008000kHz 0 +
1200000kHz 0 +
1248000kHz 0 +
1296000kHz 0 +
1416000kHz 0 +
1512000kHz 1200 = 1200 /* within the statistic time, there are 1200 system
jiffies in total, 1512M running for 1200 jiffies */
/* from below, we can see the loading status of each cpu, including user mode,
kernel mode, interrupt and idle time */
cpu0: User 0 + Nice 0 + Sys 1 + Idle 294 + IOW 0 + IRQ 0 + SIRQ 0 = 295
cpu1: User 299 + Nice 0 + Sys 1 + Idle 0 + IOW 0 + IRQ 0 + SIRQ 0 = 300
cpu2: User 1 + Nice 0 + Sys 1 + Idle 296 + IOW 0 + IRQ 0 + SIRQ 0 = 298
cpu3: User 300 + Nice 0 + Sys 1 + Idle 0 + IOW 0 + IRQ 0 + SIRQ 0 = 301

/* check the ratio of running time for each frequency through cpufreq node, time
unit: jiffies */
cat /sys/devices/system/cpu/cpu0/cpufreq/stats/time_in_state
408000 718186
600000 548
816000 368
1008000 1578
1200000 1104
1248000 84
1296000 101
1416000 678
1512000 47495

/* check the interrupt quantity of all peripherals */
cat /proc/interrupts

```

	CPU0	CPU1	CPU2	CPU3			
1:	0	0	0	0	GICv2	29	Edge
arch_timer							
2:	181898	165057	636772	839244	GICv2	30	Edge
arch_timer							
5:	180743	39000	28905	65189	GICv2	62	Level rk_timer
13:	260634	0	0	0	GICv2	39	Level
ff180000.i2c							
14:	354805	0	0	0	GICv2	40	Level
ff190000.i2c							
15:	0	0	0	0	GICv2	41	Level
ff1a0000.i2c							
...							

3.3.2 VDD_GPU

The power consumption of VDD_GPU mainly confirms if the the frequency voltage table is normal or not, if the measured voltage is consistent with the set voltage or not, using devfreq node.

Note: some chips' GPU module doesn't have separate VD and it will put GPU in VDD_LOGIC, here is needed to confirm if the voltage of VDD_LOGIC is normal or not.

```
/* acquire the frequency voltage table */
cat /sys/kernel/debug/opp/opp_summary
device          rate(Hz)    target(uV)    min(uV)    max(uV)
-----
...
platform-ff400000.gpu
                200000000    950000       950000     950000
                300000000    950000       950000     950000
                400000000    1025000      1025000     1025000
                480000000    1100000      1100000     1100000
                520000000    1150000      1150000     1150000
...

/* check the frequency scaling strategy currently used by gpu devfreq, gpu
frequency scaling is enabled with the default simple_ondemand strategy */
cat /sys/class/devfreq/ff400000.gpu/governor
simple_ondemand
Note: ff400000 of ff400000.gpu is the address of gpu register, so the name will
be different for differnt chips.

/* Set userspace strategy to fix the frequency of gpu, then set different
frequencies, compare the set voltage with measured voltage */
echo userspace > /sys/class/devfreq/ff400000.gpu/governor
cat /sys/class/devfreq/ff400000.gpu/governor
userspace

/* check the frequency points supported by gpu devfreq */
cat /sys/class/devfreq/ff400000.gpu/available_frequencies
520000000 480000000 400000000 300000000 200000000

/* set the fixed frequency */
echo 200000000 > /sys/class/devfreq/ff400000.gpu/userspace/set_freq

/* confirm current frequency */
cat /sys/class/devfreq/ff400000.gpu/cur_freq
200000000

/* confirm current voltage, and compare with measued value */
cat /sys/kernel/debug/regulator/vdd_gpu/voltage
950000

/* check gpu loading */
cat /sys/class/devfreq/ff400000.gpu/load
0@200000000Hz
```

3.3.3 VDD_LOGIC

Generally VDD_LOGIC will contain many modules, in order to manage the power consumption conveniently, it will be divided into many PD internally. The power consumption can be analyzed mainly from the following aspects:

- Confirm the running frequency and switch status of each module.

```
cat /sys/kernel/debug/clk/clk_summary
clock          enable_cnt  prepare_cnt      rate  accuracy
phase
-----
xin24m          9          10    240000000      0 0
...
  pll_gpll          1           1 12000000000      0 0
    gpll          9          20 12000000000      0 0
      clk_sdio_div50  1           1 1000000000      0 0
        clk_sdio      1           5 1000000000      0 0
          sdio_sample  0           1  500000000      0 0
            sdio_drv   0           1  500000000      0
180
      clk_emmc_div50  1           1  3000000000      0 0
        clk_emmc      1           5  3000000000      0 0
          emmc_sample  0           1  1500000000      0
42
            emmc_drv   0           1  1500000000      0
180
...
```

- Confirm the switch status of each PD.

```
cat /sys/kernel/debug/pm_genpd/pm_genpd_summary
domain          status      slaves
/device
-----
pd_gpu          off
  /devices/platform/ff400000.gpu      suspended
pd_vi          off
  /devices/platform/ff4a8000.iommu      suspended
pd_vo          on
  /devices/platform/ff460f00.iommu      active
  /devices/platform/ff470f00.iommu      suspended
  /devices/platform/ff2e0000.video-phy  suspended
  /devices/platform/ff450000.dsi        active
  /devices/platform/ff460000.vop        active
  /devices/platform/ff470000.vop        suspended
  /devices/platform/ff480000.rk_rga     suspended
pd_vpu          off
  /devices/platform/ff440440.iommu      suspended
  /devices/platform/ff442800.iommu      suspended
  /devices/platform/vpu_combo          suspended
pd_mmc_nand     on
  /devices/platform/ff380000.dwmmc      unsupported
```

```

/devices/platform/ff390000.dwmnc      unsupported
/devices/platform/ff3b0000.nandc      active
pd_gmac                             off
pd_sdcard                            off
pd_usb                               on
/devices/platform/ff300000.usb        active

```

- Generally DDR module is put in VDD_LOGIC, and the power consumption of DDR module is relatively large, use the same devfreq strategy as GPU to optimize the power consumption, so need to confirm the frequency voltage table and measured voltage. DDR also has some configurations with low power consumption, such as pd_idle, sr_idle, odt switch and some other timing configurations. The debugging process is relatively complex, you need to refer to the detailed DDR document.

```

cat /sys/kernel/debug/opp/opp_summary
device          rate(Hz)    target(uV)    min(uV)    max(uV)
-----
platform-dmc
                194000000    950000       950000     950000
                328000000    950000       950000     950000
                450000000    950000       950000     950000
                528000000    975000       975000     975000
                666000000    1000000      1000000     1000000
...

/* ddr uses dmc_ondemand frequency scaling strategy by default */
cat /sys/class/devfreq/dmc/governor
dmc_ondemand

Other commands to set the frequency and voltage are the same as GPU devfreq.

```

3.3.4 VCC_DDR

VCC_DDR supplies power mainly for DDR component and DDR-IO part of SoC. The parameters affecting the power consumption of VCC_DDR include: DDR frequency, DDR loading, DDR low power consumption configuration, DDR component type and so on. Under the same condition, the power consumption of DDR components from different vendors may have big difference.

3.3.5 VCC_IO

VCC_IO supplies power mainly for IO Pad of SoC and some peripherals. The power consumption can be analyzed from the following aspects:

- Check the working status of peripheral module, if there is leakage.
- Check if IO pin status of SoC matches with the peripheral or not, for example, IO output is high, but the connected peripheral pin is low level.

3.4 Common scenario analysis

3.4.1 Static desktop

It is mainly the display module which is working, CPU, GPU, DDR should be reduced to the lowest frequency, and enter low power consumption mode. Adjust VDD_CPU, VDD_GPU, VDD_LOGIC to the lowest voltage of opp-table, confirm the status of clk_summary and pm_genpd_summary, confirm the peripheral modules (WIFI, BT, etc.) are all closed. The static desktop generally is used as the basic power consumption of other scenarios, so need to firstly optimize its power consumption to the best.

3.4.2 Video playback

It is mainly the video decoder (VPU/RKVDEC) which is working, GPU generally is closed. Especially confirm if the running frequency of DDR and voltage of VDD_LOGIC are normal or not.

3.4.3 Game

It is mainly CPU and GPU which are working. Especially analyze the loading of CPU and GPU, frequency change, the voltages of VDD_CPU and VDD_GPU are normal or not.

3.4.4 Deepsleep

Generally VDD_CPU and VDD_GPU will turn off the power supply, VDD_LOG only reserves the power supply for some resume module, so need to focus on the power consumption analysis of IO, DDR components and some peripherals.

4. Power consumption optimization strategy

4.1 CPU optimization

- Adjust cpufreq parameter.

```
/* the default frequency scaling strategy used is interactive, relative
parameters are as follows: */
ls -l /sys/devices/system/cpu/cpu0/cpufreq/interactive
go_hispeed_load      /* when the loading is larger than go_hispeed_load and the
frequency is smaller than hispeed_freq, directly jump to hispeed_freq */
hispeed_freq         /* when jumping from low frequency to high frequency, need to
jump to hispedd_freq first */
above_hispeed_delay /* when the frequency is larger than hispeed_freq, the time
duration before each frequency increase */
```



```

min_sample_time    /* after each frequency increase, if it is to reduce the
frequency next time , the time duration before frequency reduce */
target_loads       /* the target loading of the frequency scaling */
timer_rate         /* the loading sampling time, unit:us */
timer_slack        /* the loading sampling time after cpu entering idle */
boost              /* when the frequency is smaller than hispeed_freq, keep
boost to hispeed_freq */
boostpulse         /* when the frequency is smaller than hispeed_freq, boost to
hispeed_freq, keep a while */
boostpulse_duration /* time duration of boostpulse, unit:us */
io_is_busy         /* whether to compute io wait to cpu loading */

```

We mainly adjust three parameters: hispeed_freq, target_loads, timer_rate:

1. hispeed_freq: select an appropriate transition frequency, to make cpu stable in the medium frequency, with the best power consumption, too large or too small will cause cpu jump to high frequency easily and increase the power consumption.
2. target_loads: easier to run with low frequency after this value is increased, both the power consumption and the performance will be reduced.
3. timer_rate: easier to run with low frequency after this value is increased, both the power consumption and the performance will be reduced.

- Close some cpu, limit the highest frequency of cpu.

```

/* close cpu2, cpu3 */
echo 0 > /sys/devices/system/cpu/cpu2/online
echo 0 > /sys/devices/system/cpu/cpu3/online

/* set the max frequency of cpu0 to 1200MHz */
echo 1200000 > /sys/devices/system/cpu/cpu0/cpufreq/scaling_max_freq

```

- SoC with ARM Big-Little architecture can bind the tasks with high loading to little cores through CPUSSET since the energy efficiency of the little core is better.

/* Note: SoC with SMP architecture can also bind the tasks to some cpu so that other cpus can enter low power consumption mode, but maybe it will make cpu easy to run with high frequency, which will increase the power consumption. */

```

/* create group of litte core*/
mkdir /dev/cpuset/little

/* set cpu used by group of little core */
echo 0-3 > /dev/cpuset/little/cpus

/* add pid=1111 task into the group of little core */
echo 1111 > /dev/cpuset/little/tasks

/* Android system creates several groups by default, the framework layer puts the
tasks into differcnt groups, you can adjust cpus of each group, analyze the
power consumption */
ls /dev/cpuset
background
foreground
system-background

```

top-app

- Limit the cpu bandwidth of the tasks with high loading through CPUCTL (need to enable the macro CONFIG_CFS_BANDWIDTH).

```
/* create the group of bandwidth limitation */
mkdir /dev/cpuctl/mygroup

/* set the cycle of bandwidth limitation as 10ms */
echo 10000 > /dev/cpuctl/mygroup/cpu.cfs_quota_us

/* within each cycle, total running time of the tasks in the group cannot exceed
5ms, this value can be larger than cfs_quota_us, because it is the total running
time of multiple cpus */
echo 5000 > /dev/cpuctl/mygroup/cpu.cfs_period_us

/* add relative tasks into the group */
echo 1111 > /dev/cpuctl/mygroup/tasks
echo 1112 > /dev/cpuctl/mygroup/tasks

/* cpu.shares means to limit the bandwidth of the task through weight, used for
performance optimization, without affecting the power consumption */
/dev/cpuctl/mygroup/cpu.shares
```

4.2 DDR optimization

- Frequency scaling with scenario: configure different DDR frequencies for different scenarios, such as 4K video, video recording, dual display and so on.

```
/* scenario definition */
include/dt-bindings/clock/rk_system_status.h
#define SYS_STATUS_NORMAL      (1<<0)
#define SYS_STATUS_SUSPEND     (1<<1)
#define SYS_STATUS_IDLE        (1<<2)
#define SYS_STATUS_REBOOT      (1<<3)
#define SYS_STATUS_VIDEO_4K     (1<<4)
#define SYS_STATUS_VIDEO_1080P (1<<5)
...

/* configure the frequencies for different scenarios in dts */
arch/arm64/boot/dts/rockchip/px30.dtsi
dmc: dmc {
    compatible = "rockchip,px30-dmc";
    ...
    system-status-freq = <
        /*system status      freq(KHz)*/
        SYS_STATUS_NORMAL    528000
        SYS_STATUS_REBOOT     450000
        SYS_STATUS_SUSPEND    194000
        SYS_STATUS_VIDEO_1080P 450000
        SYS_STATUS_BOOST      528000
        SYS_STATUS_ISP        666000
    >
```

```

SYS_STATUS_PERFORMANCE 666000
>;
...

/* acquire the current scenario */
cat /sys/class/devfreq/dmc/system_status
0x401

```

- Frequency scaling with loading: monitor the loading, automatically adjust DDR frequency, frequency scaling with loading may cause the reduction of the performance, you can fix DDR frequency in some scenario considering the frequency scaling with scenario.

```

/* configure the parameter of frequency scaling with loading in dts, need to open
dfi node to monitor DDR utility ratio */
dmc: dmc {
    compatible = "rockchip,px30-dmc";
    ...

    /* use dfi to monitor the utility ratio of DDR */
    devfreq-events = <&dfi>;

    /*
     * the threshold of frequency scaling:
     * when the utility ratio is over 40%, adjust to the highest frequency.
     * when the loading is less than 40% and larger than 40%-20%, maintain
current frequency.
     * when the loading is less than 40%-20%, it will adjust the frequency to a
certain value to make the loading to be around 40%-20%/2.
     */
    upthreshold = <40>;
    downdifferential = <20>;

    /* check the DDR loading of current system */
    cat /sys/class/devfreq/dmc/load
33@528000000Hz

```

- For more detailed configuration and optimization of DDR DEVFREQ, please refer to the document 《Rockchip-Developer-Guide-Linux4.4-Devfreq》.

4.3 Thermal control optimization

When the temperature is increasing to certain degree, the power consumption will increase dramatically, especially in the case with high voltage.

- Improve the heat dissipation of hardware.
- Optimize the software thermal control strategy to avoid the big temperature fluctuation.
- Avoid the high voltage occurring in the case with high temperature through software limitation.

```
&cpu0_opp_table {
    /* when the temperature is over 85 degree, limit the max voltage of cpu to
    1.1V */
    rockchip,high-temp = <85000>;
    rockchip,high-temp-max-volt = <1100000>;

    /* or directly limit the max frequency to avoid the high voltage */
    rockchip,high-temp-max-freq = <1008000>;
};
```

4.4 Power optimization

- In voltage conversion circuit, when the voltage reduction and current are relatively large, it is recommended to use DCDC to improve the efficiency and reduce the power consumption.

For example:

Input 3.3V, output 1.0V-50mA

Power Type	Input Current	Power Consumption
LDO	50mA	165mW
DCDC(with 80% efficiency)	18.9mA	62.4mW